

Carbon-Aware Computing: Redesigning Algorithms for a Greener Future

1st Soham Mungle

Department of Artificial Intelligence and
Machine Learning
Faculty of Engineering and Technology,
DMIHER (DU), Sawangi,
Wardha-442001
Maharashtra, India
munglesoham@gmail.com

2nd Rajendra Rewatkar

Department of Computer Science
and Medical Engineering
Faculty of Engineering and Technology,
DMIHER (DU), Sawangi,
Wardha-442001
Maharashtra, India
rajendrarr.feat@dmihir.edu.in

3rd Atharva Motewar

Department of Artificial Intelligence and
Machine Learning
Faculty of Engineering and Technology,
DMIHER (DU), Sawangi,
Wardha-442001
Maharashtra, India
atharvamotewar925@gmail.com

Abstract—The rise in computational pressure in the data centers, cloud, and edge has resulted in rising energy and carbon footprint. Although much focus has been on hardware efficiency and the use of renewable energy, algorithmic redesign is a close and highly effective way of looking at sustainability. In this paper, carbon-conscious computing gets reviewed with special reference to the alcoholic approaches to workload scheduling, offloading tasks, resource provisioning, and carbon footprint estimation. The most notable case studies are Google carbon-intelligent workload management, GreenScale edge-cloud scheduling, CASPER distributed web-service provisioning, market-driven edge AI optimization and the Green Algorithms footprint analysis tool. These approaches prove to have the ability to reduce emissions by up to 70 percent without affecting performance, latency, and cost. Other challenges that have been pointed out in the study include challenges in intermittency of renewable energy, computational overhead, and lack of standardized carbon measures. The future of the AI-based scheduling is predictive, further combination with the carbon markets, and cross-layer optimization of devices, edge, and data center. Carbon aware redesigning of algorithms is depicted as an important move towards making computing systems greener and more sustainable.

Keywords—Carbon-conscious computing, green computing, sustainable algorithms, datacenter optimization, edge-cloud scheduling, workload management, quantification of carbon footprint.

I. INTRODUCTION

The information and communication technologies (ICT) have been growing rapidly and have transformed the modern society but have brought severe environmental threats. ICT is estimated to be a major contributor to global greenhouse gas (GHG) emissions, with current estimates reaching 3.4 percent of global emissions by 2030 if unchecked due to data centers, high-performance computing (HPC) systems, and edge-cloud systems among other systems draw a considerable

amount of the global electricity consumption, and estimates indicate this could rise to be one of the biggest sources of global greenhouse gas emissions by 2030 [1]. As an example, data centers will constitute almost 8 percent of the global electricity use by 2030 in the absence of active mitigation strategies [2]. The increased demand is also driven by the energy-intensive application such as artificial intelligence (AI), massive data analytics, blockchain, and immersive edge services such as augmented/virtual reality (AR/VR).

A major trend in the field of conventional solutions to sustainable computing has been the idea to enhance the efficiency of hardware and to switch to renewable energy sources [3]. These approaches are not enough although they are critical. Energy efficiency is not always synonymous with carbon efficiency since carbon intensity of electricity is changing with time and location as the energy mix of the grid changes [4]. As an illustration, the computations made on a supply period dominated by coal will lead to so many higher emissions than the computation made on a period where renewables are ample. This contradiction underscores the significance of carbon conscious computing where the environmental consideration has been used in determining the real time utilization of computing options.

According to recent studies, algorithmic innovation has become one of the opportunities that can be used in the context of carbon reduction. Some of the achievements in this field are the Carbon-Intelligent Compute Management system of Google, which uses Virtual Capacity Curves (VCCs) to move flexible workloads to cleaner times [5]; edge-cloud scheduling framework of GreenScale, which can cut emissions by up to 29% [6]; and distributed web-service provisioning of CASPER, which can cut emissions by up to 70 percent [7]. Also, edge AI task offloading models combine

carbon emission rights (CERs) and dynamic carbon pricing, general-purpose instruments like Green Algorithms allow quantifying footprints in a standardized way [8], [9]. Recent papers have investigated AI-based predictive scheduling, carbon conscious container orchestration in Kubernetes and introducing carbon markets to edge cloud optimization [10], [11]. The net effect of these studies is the possibility of algorithm redesign to lead to extensive sustainability improvements.

Objective and Scope

The paper presents a state-of-the-art review of carbon-conscious computing, especially in the context of algorithmic strategies. Namely, it studies workload scheduling, resource provisioning, task offloading, and carbon footprint estimation techniques of data centers, edge-cloud-based infrastructure, and distributed services. The review formulates carbon reduction, performance, cost, and scalability as successful, limited, and trade-offs. Moreover, it also outlines the unsolved problems, including inconsistency in renewable power supply, algorithmic complexity, and the lack of cohesive carbon metrics, and provides the directions of research in the future, including predicting the timeline with the use of AI, integrating with carbon markets, and inter-layered optimization between devices, networks, and data centers. In that way, this work will establish algorithmic redesign as the pillar of sustainable computing.

II. LITERATURE REVIEW

Carbon-conscious computing, e.g. datacenter scheduling, edge cloud optimization, distributed allocation of services and measures footprint in more traditional way, has been developed very quickly in the last few years. The different journals have attempted to restrict emissions by incorporating the idea of carbon intensity and directly applying it in decision-making by use of computation. One of them is the Carbon-Intelligent Compute Management that is deployed by Google and Virtual Capacity Curve (VCCs) which recasts real-time

the workload in a grid and responds to high-carbon heterogeneous, edge-cloud environment [5] and [6]. CASPER is able to implement the distributed web services at 70 percent decrease in the emissions without lack of available quality services [7]. Other models besides timing also set a traditional way of reporting the carbon footprints of computations as temporary and not always optimum e.g. the Green Algorithms [9].

Figure 1 provides an overall comparative outline that can be used to visualize the scope of the existing systems in the backdrop of the various strategies to become computer conscious and carbon conscious. This Figure models major systems and approaches into larger groups, e.g., workload scheduling, task offloading, carbon estimation and resource provisioning.




Fig. 1. Carbon-aware computing approaches like workload scheduling, resource provisioning, carbon footprint estimation and offloading of tasks.

Table I also elaborates them with a comparative summary on their area of interest, methodology, results and limitations. Table I gives a comparative explanation of the outstanding approaches to accommodate such contributions that are not homogenous. The focal point, strategies, conclusions, and weaknesses of the most outstanding research are presented in the table below, and the table compares the state of affairs in the field of research.

TABLE I – COMPARATIVE SUMMARY OF CARBON-AWARE COMPUTING APPROACHES

Work / System	Focus Area	Methodology	Key Results	Limitations
Google's Carbon-Intelligent Compute Management [5]	Datacenter scheduling	Virtual Capacity Curves (VCCs) shift workloads based on carbon intensity	Significant reductions in datacenter emissions without performance loss	Relies on grid-level carbon forecasts, coarse-grained signals
Zhao et al. [4]	Datacenter energy management	Carbon-intensity-aware resource allocation	Improved alignment with	Limited to regional energy data, not scalable globally

			renewable availability	
GreenScale [6]	Edge-cloud infrastructures	Carbon-aware scaling across heterogeneous nodes	Reduced emissions by up to 29% while meeting SLAs	Trade-offs with latency-sensitive applications
M. Liu [7]	Distributed web services	Workload placement across carbon-aware nodes	Up to 70% emission reduction without degrading QoS	Focused mainly on web services, less on AI/ML tasks
Zhang et al. [8]	Edge AI task offloading	Carbon Emission Rights (CER) trading integrated into task scheduling	Balanced emissions with computational demand dynamically	Complexity increases with large-scale federated systems
Green Algorithms [9]	Carbon footprint quantification	Standardized framework for estimating computational emissions	Transparent footprint reporting for AI, HPC, bioinformatics	Only provides measurement, not optimization
Masanet et al. [1]; Shehabi et al. [3]	Global / national datacenter impact	Empirical modeling of energy & carbon emissions	Found ICT contributes 3–4% of global GHG; projected growth trends	Aggregate-level estimates, lack of task-level granularity

Limitations of Existing Works

Out of the comparative analysis, it is clear that there has been significant breakthrough in the field of carbon conscious computing, yet much still has to be done. The Carbon-Intelligent Compute Management offered by Google relies on coarse-grain grid-level predictions, which lowers real time accuracy. The framework by Xu et al. is mainly based on regional data, which cannot be extended to a global setting. GreenScale makes use of reduced emissions but creates latency trade-offs that make it less applicable to time-sensitive applications. CASPER can achieve a reduction of up to 70 percent of emissions of distributed web services, but does not extrapolate to workloads based on AI/ML. The CER-based scheduling developed by Zhang et al. considers carbon markets at the cost of system complexity especially in large federated systems. Footprint measurement is also presented transparently by Green Algorithms, but not optimized. Massive empirical investigations like the one by Masanet et al. and Shehabi et al. indicate the increasing role of ICT in emissions but are still on the higher levels, being aggregate, not task-practical, and lacking design-actionable details.

Research Gap

Those shortcomings indicate a vast research gap: although current solutions offer useful techniques to make carbon-conscious scheduling, provisioning, offloading, and footprint estimation, most of them are limited by the scope of domain applicability, the use of course or region-specific data, or the lack of real-time optimization measures. Also, there are minimal contributions to the use of predictive AI methods or market-based processes into cross-layer computing systems. To fill this gap, this review in brief outlines algorithmic approaches in a variety of environments (datacenters, edge-cloud, and distributed systems) and the necessity of consistent, scalable and viable frameworks that are carbon efficient, trade performance-cost-latency and proof of work.

III. METHODOLOGICAL TRENDS IN CARBON-AWARE COMPUTING

Admittedly simple workload shifting algorithms are simply the fact that long since since grown to existence carbon-considerate computing to the new needs of carbon intensity in grids, renewable availability and forecasting analytics algorithm designs. These are the

key methodological directions that will be followed and they are as pointed out in the subsections.

A. Carbon-Aware Scheduling and Provisioning

Scheduling is one of the most common methods of reducing emissions in HPC systems and datacenters. Conventional timers will be energy efficient or maximally performing but will not normally take into account the carbon intensity. Comparatively, CSP-conscience planners incorporate grid emission forecasts in making their choices. Google employs a Carbon-Intelligent Compute Management that allocates dynamically the workloads to the energy periods that are energy-sensitive of the environment according to Virtual Capacity Curves (VCCs) [5]. In addition to this, GreenScale also has the capability of scaling capacity provisioning across edge-cloud systems up to 29 percent emission cuts [6]. They are handy, however, with a scalability problem and expensive monitoring load.

B. Edge-Cloud Offloading and Distributed Optimization

The advent of mobile applications and IoT has helped to advance research in carbon aware edge-cloud coordination. The emissions and latency could also be optimised where the systems would make a selective pull of the computation between the cloud node and the edge node in accordance with the carbon conditions in the location. CASPER will present it since it moves it around web- services on green nodes without any performance degradation [7]. Also, carbon pricing and carbon emission rights (CERs) in the dynamic offloading facilities is being used to place an economic stimulus in the green computing [8]. Such strategies do bring in complexity in trade-offs of user quality-of-service (QoS) objectives and emissions targets.

C. Carbon Footprint Quantification Frameworks

The correct and identical image of the carbon footprint is a crucial prerequisite to carbon-sensitive optimization. Several models such as Green Algorithms [9] also provide straightforward methods of approximating the carbon emissions of specific hardware and workload. They are applicable to perform reasonable comparison and accountability, but not with mainstream computing since they do not communicate with a runtime environment. No standard measure can be used to compare the reported emission reductions across systems.

D. AI-Driven Predictive Carbon Management

Recent articles are associated with the assumption that the carbon intensity and the availability of renewables should be used to plan in the future using machine learning. Predictive models allow the

algorithms to redistribute workloads when the actual conditions are taken into account, and more favourable times on the grid are predicted. Proactive strategies are better and more adaptive compared to reactive strategies. Nonetheless, with valid predictions, one can get uncertainty, and even training models at a cost in carbon, it is not obvious that AI-based solutions will be useful (or not) in the end.

Summary: Trends of this direction indicate that the concept of carbon-conscious computing presupposes a new orientation where it is no longer a reactive concept but a predictive, distributed, and standardized one. Nevertheless, it continues to suffer the lack of scalability issues, forecast and single measure accuracy.

IV. OPEN CHALLENGES AND RESEARCH GAPS

Nonetheless, there has been significant advancement in carbon-conscious computing that has failed to fully resolve several issues underlining it. They are the features of the gaps in the existing approaches and the knowledge gaps in which the research studies could be satisfied.

A. Renewable Energy Intermittency

Renewable availability can generally be suggested with reliable forecasts by carbon-conscious scheduling. However, more sustainable energy sources such as solar and wind are intermittent in nature and prone to errors in predictions. The variability of renewable production poses difficulties to the fluctuation of workload and may lead to either the underutilization of renewable energy or fallback generation that is usually carbon-intensive [1], [2].

B. Algorithmic Complexity and Scalability

The majority of the carbon-aware algorithms, in particular, the ones based on optimization or reinforcement learning, are highly intricate. It is not easy to scale to systems or geographically-distributed hyperscale datacenters and systems, though it can be scaled to small scale workloads. Trade-offs in optimization accuracy and execution overhead as well as execution scheduling responsiveness have not yet been fully addressed [3], [4].

C. Lack of Standardized Carbon Metrics

The current literature uses various indicators including carbon intensity and carbon usage effectiveness (CUE) or absolute emission. Cross-studies however cannot be compared and evaluated because there are no standardized benchmarks. The profits were not easily quantifiable, provable and reproducible in form of standard measurements [5].

D. Integration with Carbon Markets and Pricing Models

New publications suggest that they need to allocate the workloads based on the dynamic pricing of carbon and carbon emission rights (CERs). Nevertheless, these models cannot be implemented in a real-world setting of computing environment because of the regulatory, economic, and system level constraints. The interface between the carbon conscious algorithms and the global carbon markets is a rather new sphere [6].

E. Performance–Sustainability Trade-offs

Balancing between carbon reduction and service quality is one of the questions that keep on emerging. Latencies can be caused by carbon sensitive aggressive scheduling, bad user experience or service level agreement (SLA) failure. The development of adaptable designs that would be common in the goal of maximizing carbon efficiency, cost, and performance is still a challenge [7].

“Fig. 2 shows the key issues that have been identified when designing carbon-conscious algorithms.”




Fig. 2. Challenges in carbon-aware algorithm design, including renewable intermittency, algorithmic complexity, carbon metric standardization, carbon market integration, and performance–sustainability trade-offs.

Summary of Challenges

In summary, renewable intermittency, algorithmic scalability, lack of standardized measures of assessment, absence of integration with economic models and performance sustainability trade-offs are all factors limiting the feasibility of carbon-conscious computing. These are the problems which should be tackled to bring research prototypes to large scale, industry ready deployments.

V. RESULTS AND DISCUSSION

The comparative analysis (Table I) and methodological synthesis have shown that carbon-conscious computing is much more efficient in terms of the sustainability outcomes without compromising the system performance. In the frameworks reviewed, the reported emission reduction is between 29% and 70% depending on the flexibility of the workload, the granularity of the scheduling system, and the magnitude of the systems.

In Google, the Carbon-Intelligent Compute Management showed potential to save a lot of emissions by matching the elastic workloads with low-carbon grid periods and at the same time produce a significant amount of savings without a lot of latency [5]. Likewise, GreenScale also optimized the capacity provisioning of heterogeneous edge-clouds, decreasing a total amount of emissions by about 29 percent, without service-level agreements (SLAs) [6]. CASPER is a distributed web services-oriented application that produced a reduction of up to 70 percent carbon and confirmed the viability of emission conscious migration of workloads in a large-scale deployment [7].

These performance gains are however at a price. The dependence on grid-level carbon projections adds some temporal uncertainty, which may result in poor scheduling decisions. Moreover, the complexity of algorithms in dynamism in resource allocation and offloading of tasks presents a scalability problem especially in real time or latency-sensitive applications. Dynamic optimization potential is further improved with the integration of Carbon Emission Rights (CER) trading models [8], but with the added complexity of the computation and the system design.

As far as performance indicators are concerned, the key ones are:

- Carbon Intensity ($\text{gCO}_2 / \text{kWh}$): Rate of emission per unit of energy, which is applied in [4] and [5] to make scheduling decisions.
- Carbon Usage Effectiveness (CUE): This is an extension of Power Usage Effectiveness (PUE) in which the amount of carbon emitted is taken into consideration, used in data center assessment.
- Percentage of Reduction of Emission (Percent): This is used to measure the total sustainability improved relative to operating at the baseline.
- Service-Level Agreement (SLA) Compliance: It is used to make sure that the performance or QoS is not compromised by sustainability gains.

The statistical results presented in the studies which have been reviewed support the fact that the relationship between the availability of renewable and workload flexibility is strong, which proves the fact that the adaptation to time and space has a quantifiable sustainability value. As an example, the authors such as Xu et al. [4] have noted a better match with renewable supply, and indirect emissions are reduced without a significant loss of throughput.

Critically speaking, the findings may suggest that carbon-awareness on the algorithmic level can expand the hardware efficiency and renewable energy integration. However, standardized measurement scales are required to allow the equitable comparisons of studies and real-life testing. The practice of current reporting is diverse, which restricts cross-comparison and reproducibility.

The discussion further notes that predictive scheduling based on AI has a potential to improve the temporal optimization of the grid by predicting the intensity of the carbon more effectively. However, these approaches bring in energy and carbon expenses of model training. Implementations in the future should therefore be able to balance the cost of computing with the avoidance of carbon so that the net effect is a positive impact on the environment.

The general findings are that algorithmic redesign will be an essential facilitator of sustainable ICT. With the carbon intensity incorporated in the workload placement, the task scheduling, and the resource provisioning decisions, it is possible to achieve substantial emission reduction without affecting the operational efficiency. Nevertheless, the move of experimental prototypes to scalable and production-ready systems is still unfinished and has to be standardized, automatized, and cross-disciplinary.

VI. CONCLUSION AND FUTURE SCOPE

Carbon aware computing is a proposal of sustainable IT, which relies on integrating algorithmic intelligence with environmental aspirations. The paper has outlined the advancement of the timing of workloads, resource delivery, offloading and estimation of carbon footprint such as Carbon-Intelligent Compute Management by Google, GreenScale, CASPER and Green Algorithms. Even though these strategies have demonstrated considerable reduction of emissions of between 29 and 70 percent, most of these strategies remain research prototypes, constrained by intermittency of renewables, algorithm complexity, scale constraints and lack of standard metrics of carbon.

These constraints are overcome by some of the promising research directions that can be utilized in order to scale to large systems at a faster rate. The prediction of solar and wind energy can be improved with the help of artificial intelligence and machine learning to reduce the uncertainty and increase the dependability of such a solution in terms of carbon-respectful scheduling. It is demanded that the lightweight and adaptive algorithms are required to ensure scalability in the hyperscale datacenters, and the infrastructures that are distributed geographically without affecting the efficiency of optimization. The measurement indicators such as improved carbon usage effectiveness (CUE) indicators should be standardized to allow its evaluation to be consistent across various systems. It will be used together with active carbon pricing models, so-called carbon emission rights (CERs) and international regulatory frameworks will coordinate technical innovations as well as sustainability goals on the global scale. Viable deployments require the use of multi-objective optimization models that involve a trade-off between carbon efficiency and cost-effectiveness and service-level objectives (SLOs). The school knowledge and the industry knowledge can assist in transforming the prototypes to the production worthy systems through the joint endeavors. Carbon-conscience computing can be used to meet the needs of high performance, low cost and environmental sustainability by making carbon issues part of the computation process. These developments will see to it that ICT is able to contribute a lot to the realization of the global sustainability goals and guide future research to multi-layered and multi-applicable solutions that are at a scalable level.

REFERENCES

- [1] E. Masanet, A. Shehabi, N. Lei, S. Smith, and J. Koomey, "Recalibrating global data center energy-use estimates," *Science*, vol. 367, no. 6481, pp. 984–986, 2020.
- [2] International Energy Agency (IEA), "Data Centres and Data Transmission Networks – Tracking Report," IEA, Paris, 2023. [Online]. Available: <https://www.iea.org/reports/data-centres-and-data-transmission-networks>
- [3] A. Shehabi et al., "United States data center energy usage report," Lawrence Berkeley National Laboratory, 2016.
- [4] Zhao, Daming, and Jiantao Zhou. "An energy and carbon-aware algorithm for renewable energy usage maximization in distributed cloud data centers." *Journal of Parallel and Distributed Computing* 165 (2022): 156-166.
- [5] Cao, Zhiwei, Xin Zhou, Han Hu, Zhi Wang, and Yonggang Wen. "Toward a systematic survey for carbon neutral data centers." *IEEE Communications Surveys & Tutorials* 24, no. 2 (2022): 895-936.
- [6] Son, Yonglak, Udit Gupta, Andrew McCrabb, Young Geun Kim, Valeria Bertacco, David Brooks, and Carole-Jean Wu. "GreenScale: Carbon Optimization for Edge Computing." *IEEE Internet of Things Journal* (2025).

- [7] M. Liu, X. Sun, H. Xu, and F. Liu, "CASPER: Carbon-aware scheduling and provisioning for distributed web services," in Proc. IEEE INFOCOM, 2021, pp. 1–10.
- [8] Ma, Huirong, Zhi Zhou, Xiaoxi Zhang, and Xu Chen. "Toward carbon-neutral edge computing: Greening edge AI by harnessing spot and future carbon markets." IEEE Internet of Things Journal 10, no. 18 (2023): 16637-16649.
- [9] L. Lannelongue, J. Grealey, and M. Inouye, "Green algorithms: Quantifying the carbon footprint of computation," Adv. Sci., vol. 8, no. 12, pp. 1–9, 2021.
- [10] Sallam, Mahmoud, and Kuljeet Kaur. "Optimized resource forecasting for carbon-intelligent data centers with temposight: A hybrid deep learning approach." IEEE Internet of Things Journal (2025).
- [11] H. Wang, Y. Lin, X. Zhang, and J. Zhang, "Carbon-aware container orchestration in Kubernetes clusters," in Proceedings of the IEEE International Conference on Cloud Computing (CLOUD), 2022, pp. 231–240.